

```
In [ ]: # Initialize Otter
import otter
grader = otter.Notebook("lab03.ipynb")
```

1 Lab 3: Data Cleaning and EDA

In this lab you will be working on visualizing a dataset from the City of Berkeley containing data on calls to the Berkeley Police Department. Information about the dataset can be found [at this link](#).

Due Date: Saturday, July 2, 11:59 PM PT.

1.0.1 Content Warning

This lab includes an analysis of crime in Berkeley. If you feel uncomfortable with this topic, **please contact your GSI or the instructors**.

1.0.2 Collaboration Policy

Data science is a collaborative activity. While you may talk with others about this assignment, we ask that you **write your solutions individually**. If you discuss the assignment with others, please **include their names** in the cell below.

Collaborators: *list names here*

1.1 Setup

In this lab, we'll perform Exploratory Data Analysis and learn some preliminary tips for working with matplotlib (a Python plotting library). Note that we configure a custom default figure size. Virtually every default aspect of matplotlib [can be customized](#).

Collaborators: *list names here*

```
In [1]: import pandas as pd
import numpy as np
```

```
import zipfile
import matplotlib
import matplotlib.pyplot as plt

plt.rcParams['figure.figsize'] = (12, 9)
```

2 Part 1: Acquire the Data

1. Obtain data To retrieve the dataset, we will use the `ds100_utils.fetch_and_cache` utility.

```
In [2]: # just run this cell
        from ds100_utils import download_lab3_data

        dest_path = download_lab3_data()
        print(f'Located at {dest_path}')
```

Using cached version that was downloaded (UTC): Sat Jun 11 02:17:51 2022
Located at data/lab03_data_sp22.zip

2. Unzip file We will now directly unzip the ZIP archive and start working with the uncompressed files.

```
In [4]: # just run this cell
        my_zip = zipfile.ZipFile(dest_path, 'r')
        my_zip.extractall('data')
```

Note: There is no single right answer regarding whether to work with compressed files in their compressed state or to uncompress them on disk permanently. For example, if you need to work with multiple tools on the same files, or write many notebooks to analyze them—and they are not too large—it may be more convenient to uncompress them once. But you may also have situations where you find it preferable to work with the compressed data directly.

Python gives you tools for both approaches, and you should know how to perform both tasks in order to choose the one that best suits the problem at hand.

3. View files

Now, we'll use the `os` package to list all files in the `data` directory. `os.walk()` recursively traverses the directory, and `os.path.join()` creates the full pathname of each file.

If you're interested in learning more, check out the Python3 documentation pages for `os.walk` ([link](#)) and `os.path` ([link](#)).

We use Python 3 [format strings](#) to nicely format the printed variables `dpath` and `fpath`.

```
In [6]: # just run this cell
import os

for root, directories, filenames in os.walk('data'):
    # first, print out all directories
    for directory in directories:
        dpath = os.path.join(root, directory)
        print(f"d {dpath}")

    # next, print out all files
    for filename in filenames:
        fpath = os.path.join(root, filename)
        print(f"  {fpath}")
```

```
d data/secret
data/Berkeley_PD_-_Calls_for_Service.csv
data/ben_kurtovic.py
data/dummy.txt
data/hello_world.py
data/lab03_data_sp22.zip
data/secret/do_not_readme.md
```

In this Lab, we'll be working with the `Berkeley_PD_-_Calls_for_Service.csv` file. Feel free to check out the other files, though.

3 Part 2: Clean and Explore the Data

Let's now load the CSV file we have into a `pandas.DataFrame` object and start exploring the data.

```
In [7]: # just run this cell
calls = pd.read_csv("data/Berkeley_PD_-_Calls_for_Service.csv")
calls.head()
```

```
Out[7]:
```

	CASENO	OFFENSE	EVENTDT	EVENTTM	\
0	21014296	THEFT MISD. (UNDER \$950)	04/01/2021 12:00:00 AM	10:58	
1	21014391	THEFT MISD. (UNDER \$950)	04/01/2021 12:00:00 AM	10:38	
2	21090494	THEFT MISD. (UNDER \$950)	04/19/2021 12:00:00 AM	12:15	
3	21090204	THEFT FELONY (OVER \$950)	02/13/2021 12:00:00 AM	17:00	
4	21090179	BURGLARY AUTO	02/08/2021 12:00:00 AM	6:20	

	CVLEGEND	CVDOW	InDbDate	\
0	LARCENY	4	06/15/2021 12:00:00 AM	
1	LARCENY	4	06/15/2021 12:00:00 AM	
2	LARCENY	1	06/15/2021 12:00:00 AM	
3	LARCENY	6	06/15/2021 12:00:00 AM	
4	BURGLARY - VEHICLE	1	06/15/2021 12:00:00 AM	

	Block_Location	BLKADDR	\
0	Berkeley, CA\n(37.869058, -122.270455)	NaN	
1	Berkeley, CA\n(37.869058, -122.270455)	NaN	
2	2100 BLOCK HASTE ST\nBerkeley, CA\n(37.864908,...	2100 BLOCK HASTE ST	
3	2600 BLOCK WARRING ST\nBerkeley, CA\n(37.86393...	2600 BLOCK WARRING ST	
4	2700 BLOCK GARBER ST\nBerkeley, CA\n(37.86066,...	2700 BLOCK GARBER ST	

	City	State
0	Berkeley	CA
1	Berkeley	CA
2	Berkeley	CA
3	Berkeley	CA
4	Berkeley	CA

We see that the fields include a case number, the offense type, the date and time of the offense, the “CVLEGEND” which appears to be related to the offense type, a “CVDOW” which has no apparent meaning, a date added to the database, and the location spread across four fields. We can read more about each field from the City of the Berkeley’s [open dataset webpage](#).

Let’s also check some basic information about this DataFrame using the `DataFrame.info` ([documentation](#)) and `DataFrame.describe` methods ([documentation](#)).

```
In [8]: # df.info() displays
        # name and type of each column,
        # number of non-null entries, and
        # size of dataframe
        calls.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2632 entries, 0 to 2631
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   CASENO           2632 non-null   int64
1   OFFENSE          2632 non-null   object
2   EVENTDT          2632 non-null   object
3   EVENTTM          2632 non-null   object
4   CVLEGEND         2632 non-null   object
5   CVDOW            2632 non-null   int64
6   InDbDate         2632 non-null   object
7   Block_Location   2632 non-null   object
8   BLKADDR          2612 non-null   object
```

```

9   City                2632 non-null   object
10  State               2632 non-null   object
dtypes: int64(2), object(9)
memory usage: 226.3+ KB

```

Note that the BLKADDR column only has 2612 non-null entries, while the other columns all have 2632 entries. This is because the `.info()` method only counts non-null entries.

```
In [9]: calls.describe()
```

```

Out[9]:
          CASENO      CVDOW
count  2.632000e+03  2632.000000
mean    2.095978e+07    3.071049
std      2.452665e+05    1.984136
min      2.005721e+07    0.000000
25%      2.100568e+07    1.000000
50%      2.101431e+07    3.000000
75%      2.102256e+07    5.000000
max      2.109066e+07    6.000000

```

Notice that the functions above reveal type information for the columns, as well as some basic statistics about the numerical columns found in the DataFrame. However, we still need more information about what each column represents. Let's explore the data further in Question 1.

Before we go over the fields to see their meanings, the cell below will verify that all the events happened in Berkeley by grouping on the `City` and `State` columns. You should see that all of our data falls into one group.

```
In [10]: calls.groupby(["City", "State"]).count()
```

```

Out[10]:
          City      State  CASENO  OFFENSE  EVENTDT  EVENTTM  CVLEGEND  CVDOW  InDbDate  \
0  Berkeley  CA         2632     2632     2632     2632     2632     2632     2632

          Block_Location  BLKADDR
0  Berkeley  CA         2632     2612

```

When we called `head()` on the DataFrame `calls`, it seemed like `OFFENSE` and `CVLEGEND` both contained information about the type of event reported. What is the difference in meaning between the two columns? One way to probe this is to look at the `value_counts` for each Series.

```
In [11]: calls['OFFENSE'].value_counts().head(10)
```

```
Out[11]: THEFT MISD. (UNDER $950)      559
         VEHICLE STOLEN                 277
         BURGLARY AUTO                 218
         THEFT FELONY (OVER $950)      215
         DISTURBANCE                  204
         BURGLARY RESIDENTIAL          178
         VANDALISM                    166
         THEFT FROM AUTO              163
         ASSAULT/BATTERY MISD.        116
         ROBBERY                      90
         Name: OFFENSE, dtype: int64
```

```
In [12]: calls['CVLEGEND'].value_counts().head(10)
```

```
Out[12]: LARCENY                        782
         MOTOR VEHICLE THEFT           277
         BURGLARY - VEHICLE            218
         DISORDERLY CONDUCT           204
         BURGLARY - RESIDENTIAL        178
         VANDALISM                    166
         LARCENY - FROM VEHICLE        163
         ASSAULT                      150
         FRAUD                        93
         ROBBERY                      90
         Name: CVLEGEND, dtype: int64
```

It seems like OFFENSE is more specific than CVLEGEND, e.g. “LARCENY” vs. “THEFT FELONY (OVER \$950)”. If you’re unfamiliar with the term, “larceny” is a legal term for theft of personal property.

To get a sense of how many subcategories there are for each OFFENSE, we will set `calls_by_cvlegend_and_offense` equal to a multi-indexed series where the data is first indexed on the CVLEGEND and then on the OFFENSE, and the data is equal to the number of offenses in the database that match the respective CVLEGEND and OFFENSE. As you can see, `calls_by_cvlegend_and_offense["LARCENY", "THEFT FROM PERSON"]` returns 8 which means there are 8 instances of larceny with offense of type “THEFT FROM PERSON” in the database.

```
In [13]: calls_by_cvlegend_and_offense = calls.groupby(["CVLEGEND", "OFFENSE"]).size()
         calls_by_cvlegend_and_offense["LARCENY", "THEFT FROM PERSON"]
```

```
Out[13]: 8
```

3.1 Question 1

In the cell below, set `answer1` equal to a list of strings corresponding to the possible values for OFFENSE when CVLEGEND is “LARCENY”. You can type the answer manually, or you can create an expression that

automatically extracts the names.

```
In [14]: answer1 = list(calls_by_cvlegend_and_offense['LARCENY'].index) # SOLUTION
```

```
In [ ]: grader.check("q1")
```

4 Part 3: Visualize the Data

4.0.1 Matplotlib demo

You've seen some `matplotlib` in this class already, but now we will explain how to work with the object-oriented plotting API mentioned in this [matplotlib.pyplot tutorial](#) useful. In `matplotlib`, plotting occurs on a set of `Axes` which are associated with a `Figure`. An analogy is that on a blank canvas (`Figure`), you choose a location to plot (`Axes`) and then fill it in (`plot`).

There are two approaches to labeling and manipulating figure contents, which we'll discuss below. Approach 1 is closest to the plotting paradigm of MATLAB, the namesake of `matplotlib`; Approach 2 is also common because many `matplotlib`-based packages (such as `Seaborn`) explicitly return the current set of axes after plotting data. Both are essentially equivalent, and at the end of this class you'll be comfortable with both.

Approach 1: `matplotlib` (or `Seaborn`) will auto-plot onto the current set of `Axes` or (if none exists) create a new figure/set of default axes. You can plot data using methods from `plt`, which is shorthand for the `matplotlib.pyplot` package. Then subsequent `plt` calls all edit the same set of default-created axes.

Approach 2:

After creating the initial plot, you can also use `plt.gca()` to explicitly get the current set of axes, and then edit those specific axes using axes methods. Note the method naming is slightly different!

As an example of the built-in plotting functionality of `pandas`, the following example uses `plot` method of the `Series` class to generate a `barh` plot type to visually display the value counts for `CVLEGEND`.

There are also many other plots that we will explore throughout the lab.

Side note: `Pandas` also offers basic functionality for plotting. For example, the `DataFrame` and `Series` classes both have a `plot` method, which uses `matplotlib` under the hood. For now we'll focus on `matplotlib` itself so you get used to the syntax, but just know that convenient `Pandas` plotting methods exist for your own future data science exploration.

Below, we show both approaches by generating a horizontal bar plot to visually display the value counts for `CVLEGEND`. See the [barh documentation](#) for more details.

```

In [20]: # DEMO CELL: assign demo to 1 or 2.
demo = ...

calls_cvlegend = calls['CVLEGEND'].value_counts()

if demo == 1:
    plt.barh(calls_cvlegend.index, calls_cvlegend) # creates figure and axes
    print(f"Demo {demo}: Using plt methods to update plot")
    plt.ylabel("Crime Category") # uses most recently plotted axes
    plt.xlabel("Number of Calls")
    plt.title("Number of Calls by Crime Type")
elif demo == 2:
    print(f"Demo {demo}: Using axes methods to update plot")
    plt.barh(calls_cvlegend.index, calls_cvlegend) # creates figure and axes
    ax = plt.gca()
    ax.set_ylabel("Crime Category")
    ax.set_xlabel("Number of Calls")
    ax.set_title("Axes methods: Number of Calls by Crime Type")
else:
    print("Error: Please assign the demo variable to 1 or 2.")

plt.show()

```

Error: Please assign the demo variable to 1 or 2.

4.0.2 An Additional Note on Plotting in Jupyter Notebooks

You may have noticed that many of our plotting code cells end with a semicolon ; or `plt.show()`. The former prevents any extra output from the last line of the cell; the latter explicitly returns (and outputs) the figure. Try adding this to your own code in the following questions!

4.1 Question 2

Now it is your turn to make a plot using `matplotlib`. Let's start by transforming the data so that it is easier to work with.

The CVDOW field isn't named helpfully and it is hard to see the meaning from the data alone. According to the website [linked](#) at the top of this notebook, CVDOW is actually indicating the day that events happened. 0->Sunday, 1->Monday ... 6->Saturday.

4.2 Question 2a

Add a new column `Day` into the `calls` dataframe that has the string weekday (eg. 'Sunday') for the corresponding value in `CVDOW`. For example, if the first 3 values of `CVDOW` are `[3, 6, 0]`, then the first 3 values of the `Day` column should be `["Wednesday", "Saturday", "Sunday"]`.

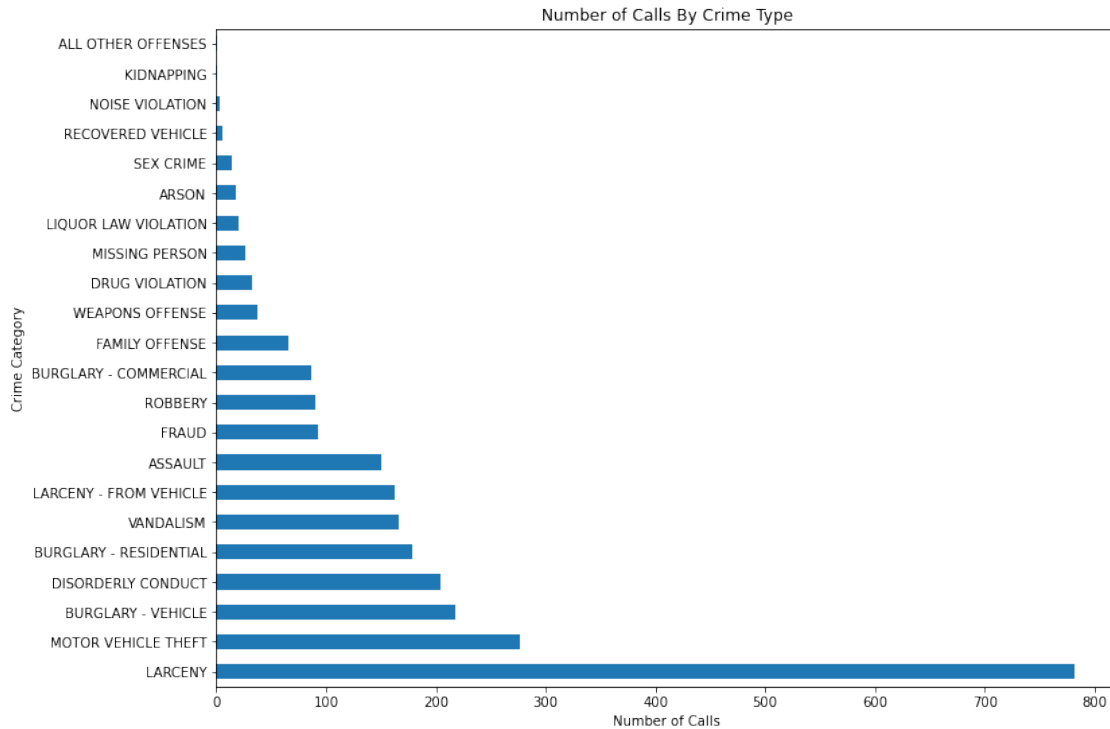
Hint: Try using the `Series.map` function on `calls["CVDOW"]`. Can you assign this to the new column `calls["Day"]`?

```
In [21]: days = ["Sunday", "Monday", "Tuesday", "Wednesday", "Thursday", "Friday", "Saturday"]
         day_indices = range(7)
         indices_to_days_dict = dict(zip(day_indices, days)) # Should look like {0:"Sunday", 1:"Monday"

         calls["Day"] = calls["CVDOW"].map(indices_to_days_dict) # SOLUTION
         # BEGIN SOLUTION NO PROMPT
         # challenge solution, commented out
         # we drop the column "Day" if it already exists, otherwise
         # duplicate "Day" columns are created
         # calls = pd.merge(calls.drop(columns="Day", errors="ignore"),
         #                  pd.DataFrame(days, columns=["Day"]),
         #                  left_on='CVDOW', right_index=True).sort_index()
         # END SOLUTION
```

```
In [ ]: grader.check("q2a")
```

```
In [24]: # just run this example cell
         ax = calls['CVLEGEND'].value_counts().plot(kind='barh')
         ax.set_ylabel("Crime Category")
         ax.set_xlabel("Number of Calls")
         ax.set_title("Number of Calls By Crime Type");
```



Challenge (OPTIONAL): You could also accomplish this part as a table left join with `pd.merge` ([documentation](#)), instead of using `Series.map`. You would need to merge `calls` with a new dataframe that just contains the days of the week. If you have time, try it out in the below cell!

```
In [25]: # scratch space for optional challenge
dow_df = pd.DataFrame(days, columns=["Day"])

...
```

Out[25]: Ellipsis

4.3 Question 2b

Now let's look at the `EVENTTM` column which indicates the time for events. Since it contains hour and minute information, let's extract the hour info and create a new column named `Hour` in the `calls` dataframe. You should save the hour as an `int`.

Hint: Your code should only require one line. **Hint 2:** The vectorized `Series.str[ind]` performs integer indexing on an array entry.

```
In [26]: calls["Hour"] = calls['EVENTTM'].str.split(':').str[0].astype(int) # SOLUTION
        calls["Hour"]
```

```
Out[26]: 0      10
         1      10
         2      12
         3      17
         4       6
         ..
        2627    12
        2628    15
        2629     0
        2630    18
        2631     2
        Name: Hour, Length: 2632, dtype: int64
```

```
In [ ]: grader.check("q2b")
```

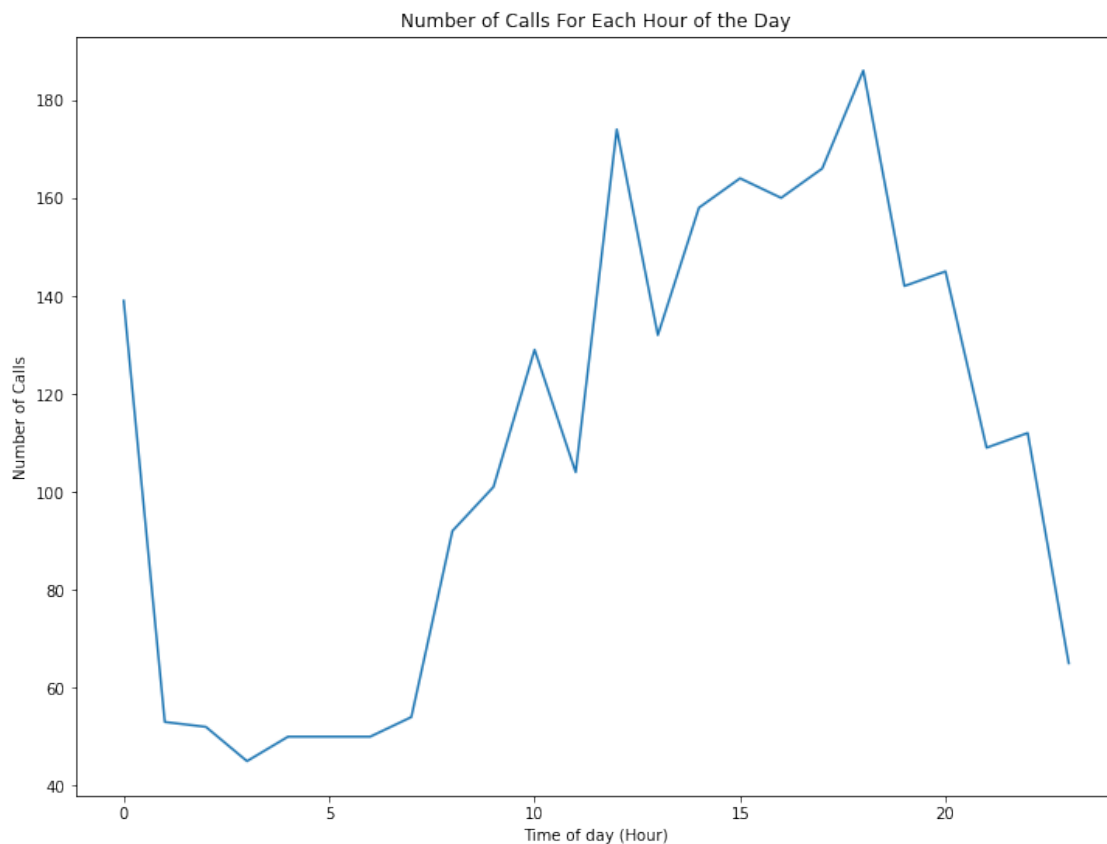
4.4 Question 2c

Using `matplotlib`, construct a line plot with the count of the number of calls (entries in the table) for each hour of the day **ordered by the time** (eg. 12:00 AM, 1:00 AM, ...). Please use the provided variable `hours` in your answer. Be sure that your axes are labeled and that your plot is titled.

Hint: Check out the `plt.plot` method in the [matplotlib tutorial](#), as well as our demo above.

```
In [29]: hours = list(range(24))
        # BEGIN SOLUTION
        calls_hour = calls["Hour"].value_counts().sort_index()
        plt.plot(calls_hour.index, calls_hour)
        ax = plt.gca()
        ax.set_xlabel("Time of day (Hour)")
        ax.set_ylabel("Number of Calls")
        ax.set_title("Number of Calls For Each Hour of the Day");
        # END SOLUTION
```

```
# Leave this for grading purposes  
ax_3d = plt.gca()
```



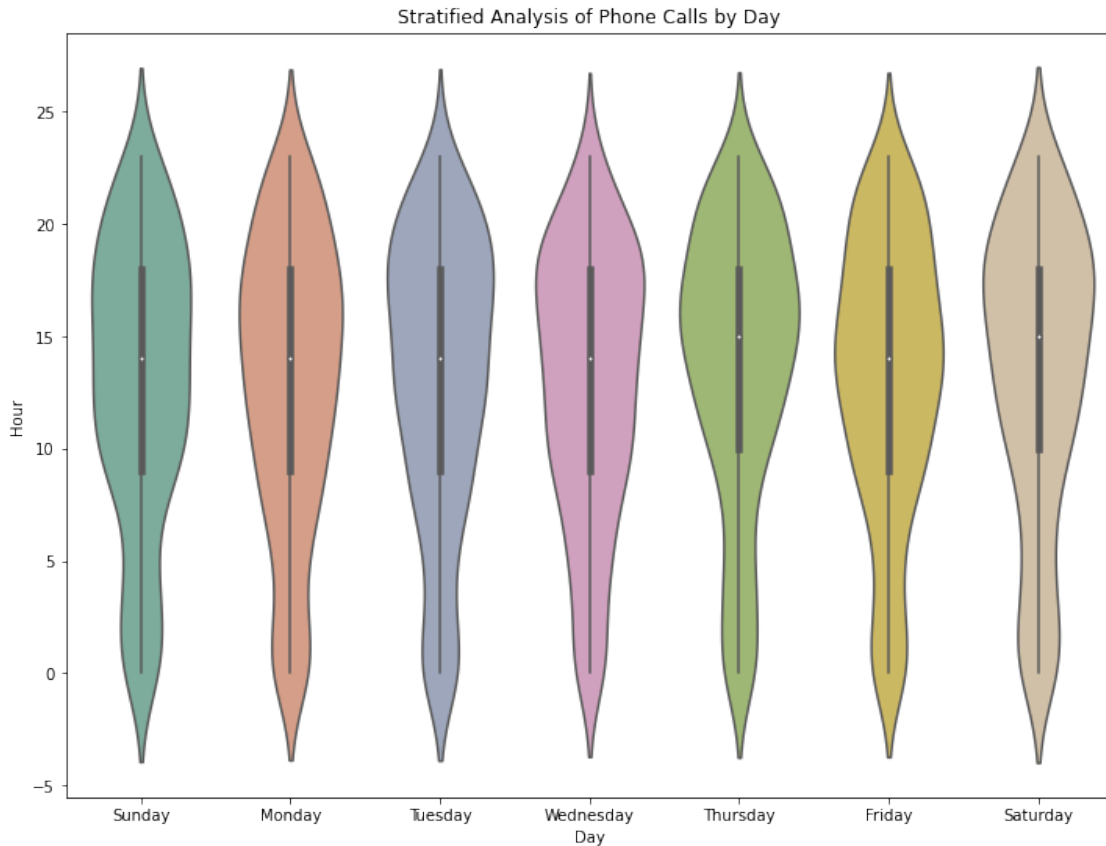
```
In [ ]: grader.check("q2c")
```

To better understand the time of day a report occurs we could **stratify the analysis by the day of the week**. To do this we will use **violin plots** (a variation of a **box plot**), which you will learn in more detail next week.

For now, just know that a violin plot shows an estimated distribution of quantitative data (e.g., distribution of calls by hour) over a categorical variable (day of the week). More calls occur in hours corresponding to the fatter part of each violin; the median hour of all calls in a particular day is marked by the white dot in the corresponding violin.

```
In [32]: # for now, just run this cell.  
         # we will learn the seaborn visualization library next week.
```

```
import seaborn as sns
ax = sns.violinplot(data=calls.sort_values("CVDOW"),
                    x="Day", y="Hour",
                    saturation=0.5, palette="Set2")
ax.set_title("Stratified Analysis of Phone Calls by Day");
```



4.5 Question 2d

Based on your line plot and our violin plot above, what observations can you make about the patterns of calls? Here are some dimensions to consider: * Are there more calls in the day or at night? * What are the most and least popular times? * Do call patterns vary by day of the week?

Type your answer here, replacing this text.

SOLUTION: In the simple line plot we see the standard pattern of limited activity early in the morning around here 6:00AM. The violin plot has no very clear patterns. However it does appear that weekends have more calls later into the night.

4.6 Question 3

In this last part of the lab, let's extract the GPS coordinates (latitude, longitude) from the `Block_Location` of each record.

```
In [33]: # an example block location entry
calls.loc[4, 'Block_Location']
```

```
Out[33]: '2700 BLOCK GARBER ST\nBerkeley, CA\n(37.86066, -122.253407)'
```

4.7 Question 3a: Regular Expressions

Use regular expressions to create a dataframe `calls_lat_lon` that has two columns titled `Lat` and `Lon`, containing the respective latitude and longitude of each record in `calls`. You should use the `Block_Location` column to extract the latitude and longitude coordinates.

Hint: Check out the `Series.str.extract` [documentation](#).

```
In [34]: calls_lat_lon = ...

# BEGIN SOLUTION NO PROMPT
# unnamed groups
calls_lat_lon = (
    calls['Block_Location']
    .str.extract("\((\d+\.\d+)\), (-\d+\.\d+)\)")
)
calls_lat_lon.columns = ['Lat', 'Lon']

# fancy version
calls_lat_lon = (
    calls['Block_Location']
    .str.extract(".*\((?P<Lat>\d*\.\d*)\), (?P<Lon>-?\d*\.\d*)\)", expand=True)
)
# END SOLUTION

calls_lat_lon.head(10)
```

```
Out[34]:
```

	Lat	Lon
0	37.869058	-122.270455
1	37.869058	-122.270455
2	37.864908	-122.267289
3	37.863934	-122.250262
4	37.86066	-122.253407
5	37.881957	-122.269551
6	37.867426	-122.269138
7	37.858116	-122.268002
8	37.868355	-122.274953
9	37.851491	-122.28563

```
In [ ]: grader.check("q3a")
```

4.8 Question 3b: Join Tables

Let's include the GPS data into our `calls` data. In the below cell, use `calls_lat_lon` to add two new columns called `Lat` and `Lon` to the `calls` dataframe.

Hint: `pd.merge` ([documentation](#)) could be useful here. Note that the order of records in `calls` and `calls_lat_lon` are the same.

```
In [37]: # BEGIN SOLUTION
# approach 1:
calls["Lat"] = calls_lat_lon["Lat"]
calls["Lon"] = calls_lat_lon["Lon"]

# approach 2:
# Remove Lat and Lon if they already existed before
calls.drop(["Lat", "Lon"], axis=1, inplace=True, errors="ignore")
# Join in the the latitude and longitude data
calls = calls.merge(calls_lat_lon, left_index=True, right_index=True)

# END SOLUTION
calls.sample(5)      # random rows
```

```
Out[37]:
```

	CASENO	OFFENSE	EVENTDT	EVENTTM	\
1615	21023792	THEFT FELONY (OVER \$950)	05/29/2021	12:00:00 AM	10:30
1722	21090048	BURGLARY AUTO	01/08/2021	12:00:00 AM	18:00
1540	21019622	VANDALISM	05/04/2021	12:00:00 AM	13:00
479	21012020	THEFT MISD. (UNDER \$950)	03/18/2021	12:00:00 AM	19:45
2220	21090621	VANDALISM	05/19/2021	12:00:00 AM	8:00

	CVLEGEND	CVDOW	InDbDate	\
1615	LARCENY	6	06/15/2021 12:00:00 AM	
1722	BURGLARY - VEHICLE	5	06/15/2021 12:00:00 AM	
1540	VANDALISM	2	06/15/2021 12:00:00 AM	
479	LARCENY	4	06/15/2021 12:00:00 AM	
2220	VANDALISM	3	06/15/2021 12:00:00 AM	

	Block_Location	\
1615	2426 MCGEE AVE\nBerkeley, CA\n(37.863593, -122...	
1722	2200 BLOCK MARIN AVE\nBerkeley, CA\n(37.891755...	
1540	1800 BLOCK 4TH ST\nBerkeley, CA\n(37.869888, -...	
479	1900 BLOCK SHATTUCK AVE\nBerkeley, CA\n(37.873...	
2220	2700 BLOCK SAN PABLO AVE\nBerkeley, CA\n(37.85...	

	BLKADDR	City	State	Day	Hour	Lat	\
1615	2426 MCGEE AVE	Berkeley	CA	Saturday	10	37.863593	
1722	2200 BLOCK MARIN AVE	Berkeley	CA	Friday	18	37.891755	
1540	1800 BLOCK 4TH ST	Berkeley	CA	Tuesday	13	37.869888	
479	1900 BLOCK SHATTUCK AVE	Berkeley	CA	Thursday	19	37.873687	
2220	2700 BLOCK SAN PABLO AVE	Berkeley	CA	Wednesday	8	37.857714	

	Lon
1615	-122.276751
1722	-122.269881
1540	-122.300618
479	-122.268616
2220	-122.288536

```
In [ ]: grader.check("q3b")
```

4.9 Question 3c: Check for Missing Values

It seems like every record has valid GPS coordinates:

```
In [41]: # just run this cell
# fraction of valid lat/lon entries
(~calls[["Lat", "Lon"]].isna()).mean()
```

```
Out[41]: Lat    1.0
Lon    1.0
dtype: float64
```


However, a closer examination of the data reveals something else. Here's the first few records of our data again:

```
In [42]: calls.head(5)
```

```
Out [42]:
```

	CASENO	OFFENSE	EVENTDT	EVENTTM	\
0	21014296	THEFT MISD. (UNDER \$950)	04/01/2021	12:00:00 AM	10:58
1	21014391	THEFT MISD. (UNDER \$950)	04/01/2021	12:00:00 AM	10:38
2	21090494	THEFT MISD. (UNDER \$950)	04/19/2021	12:00:00 AM	12:15
3	21090204	THEFT FELONY (OVER \$950)	02/13/2021	12:00:00 AM	17:00
4	21090179	BURGLARY AUTO	02/08/2021	12:00:00 AM	6:20

	CVLEGEND	CVDOW	InDbDate	\
0	LARCENY	4	06/15/2021 12:00:00 AM	
1	LARCENY	4	06/15/2021 12:00:00 AM	
2	LARCENY	1	06/15/2021 12:00:00 AM	
3	LARCENY	6	06/15/2021 12:00:00 AM	
4	BURGLARY - VEHICLE	1	06/15/2021 12:00:00 AM	

	Block_Location	BLKADDR	\
0	Berkeley, CA\n(37.869058, -122.270455)	NaN	
1	Berkeley, CA\n(37.869058, -122.270455)	NaN	
2	2100 BLOCK HASTE ST\nBerkeley, CA\n(37.864908,...	2100 BLOCK HASTE ST	
3	2600 BLOCK WARRING ST\nBerkeley, CA\n(37.86393...	2600 BLOCK WARRING ST	
4	2700 BLOCK GARBER ST\nBerkeley, CA\n(37.86066,...	2700 BLOCK GARBER ST	

	City	State	Day	Hour	Lat	Lon
0	Berkeley	CA	Thursday	10	37.869058	-122.270455
1	Berkeley	CA	Thursday	10	37.869058	-122.270455
2	Berkeley	CA	Monday	12	37.864908	-122.267289
3	Berkeley	CA	Saturday	17	37.863934	-122.250262
4	Berkeley	CA	Monday	6	37.86066	-122.253407

There is another field that tells us whether we have a valid `Block_Location` entry per record—i.e., with GPS coordinates (latitude, longitude) that match the listed block location. What is it?

In the below cell, use the field you found to create a new dataframe, `missing_lat_lon`, that contains only the rows of `calls` that have invalid latitude and longitude data. Your new dataframe should have all the same columns of `calls`.

```
In [43]: missing_lat_lon = calls[calls["BLKADDR"].isna()] # SOLUTION
missing_lat_lon.head()
```

```
Out [43]:
```

	CASENO	OFFENSE	EVENTDT	EVENTTM	\
0	21014296	THEFT MISD. (UNDER \$950)	04/01/2021	12:00:00 AM	10:58
1	21014391	THEFT MISD. (UNDER \$950)	04/01/2021	12:00:00 AM	10:38

```

215 21019124 BURGLARY RESIDENTIAL 04/30/2021 12:00:00 AM 10:00
260 21000289 VEHICLE STOLEN 01/01/2021 12:00:00 AM 12:00
633 21013362 BURGLARY AUTO 03/27/2021 12:00:00 AM 4:20

```

```

CVLEGEND CVDOW InDbDate \
0 LARCENY 4 06/15/2021 12:00:00 AM
1 LARCENY 4 06/15/2021 12:00:00 AM
215 BURGLARY - RESIDENTIAL 5 06/15/2021 12:00:00 AM
260 MOTOR VEHICLE THEFT 5 06/15/2021 12:00:00 AM
633 BURGLARY - VEHICLE 6 06/15/2021 12:00:00 AM

```

```

Block_Location BLKADDR City State Day \
0 Berkeley, CA\n(37.869058, -122.270455) NaN Berkeley CA Thursday
1 Berkeley, CA\n(37.869058, -122.270455) NaN Berkeley CA Thursday
215 Berkeley, CA\n(37.869058, -122.270455) NaN Berkeley CA Friday
260 Berkeley, CA\n(37.869058, -122.270455) NaN Berkeley CA Friday
633 Berkeley, CA\n(37.869058, -122.270455) NaN Berkeley CA Saturday

```

```

Hour Lat Lon
0 10 37.869058 -122.270455
1 10 37.869058 -122.270455
215 10 37.869058 -122.270455
260 12 37.869058 -122.270455
633 4 37.869058 -122.270455

```

```
In [ ]: grader.check("q3c")
```

4.10 Question 3d: Check Missing Values

Now let us explore if there is a pattern to which types of records have missing latitude and longitude entries.


We've implemented the plotting code for you below, but read through it and verify you understand what we're doing (we've thrown in a bonus `plt.subplots()` call, documentation [here](#)).

```

In [46]: # just run this cell
missing_by_time = (pd.to_datetime(missing_lat_lon['EVENTDT'])
                    .value_counts()
                    .sort_index()
                    )
missing_by_crime = (missing_lat_lon['CVLEGEND']
                    .value_counts()
                    / calls['CVLEGEND'].value_counts()
                    ).dropna().sort_values(ascending=False)

```

```
fig, ax = plt.subplots(2)
ax[0].bar(missing_by_time.index, missing_by_time)
ax[0].set_ylabel("Calls with Missing Data")
ax[1].barh(missing_by_crime.index, missing_by_crime)
ax[1].set_xlabel("Fraction of Missing Data per Event Type")
fig.suptitle("Characteristics of Missing Lat/Lon Data")
plt.show()
```



Based on the plots above, are there any patterns among entries that are missing latitude/longitude data? The dataset information [linked](#) at the top of this notebook may also give more context.

Type your answer here, replacing this text.

SOLUTION: While some dates have more unlabeled data than others, it seems that a small percentage of Burglary and Fraud calls don't have GPS coordinates.

4.11 Question 3d: Explore

The below cell plots a map of phonecalls by GPS coordinates (latitude, longitude); we drop missing location data.

```
In [47]: # just run this cell
import folium
import folium.plugins

SF_COORDINATES = (37.87, -122.28)
sf_map = folium.Map(location=SF_COORDINATES, zoom_start=13)
locs = calls.drop(missing_lat_lon.index)[['Lat', 'Lon']].astype('float').values
heatmap = folium.plugins.HeatMap(locs.tolist(), radius=10)
sf_map.add_child(heatmap)
```

```
Out[47]: <folium.folium.Map at 0x7f4309d406a0>
```

Based on the above map, what could be some **drawbacks** of using the location fields in this dataset to draw conclusions about crime in Berkeley? Here are some sub-questions to consider: * Is campus really the safest place to be? * Why are all the calls located on the street and often at intersections?

Type your answer here, replacing this text.

SOLUTION: This dataset is Berkeley Police crime data, not UC Police Department crime data. UC Berkeley has campus police, and that data is not included. Furthermore, calls are at intersections because the data only collects block-level granularity of locations (`BLOCKADDR` and `Block_Location`). While the location data may be useful for this type of broad human visualization, the data has missing values about a key portion of Berkeley (i.e., campus), and anyone using this dataset must acknowledge that location data granularity is block-level, and not address level.

Important: To make sure the test cases run correctly, click `Kernel>Restart & Run All` and make sure all of the test cases are still passing. Doing so will submit your code for you.

If your test cases are no longer passing after restarting, it's likely because you're missing a variable, or the modifications that you'd previously made to your DataFrame are no longer taking place (perhaps because you deleted a cell).

You may submit this assignment as many times as you'd like before the deadline.

You must restart and run all cells before submitting. Otherwise, you may pass test cases locally, but not on our servers. We will not entertain regrade requests of the form, "my code passed all of my local test cases, but failed the autograder".

4.12 Congratulations!

Congrats! You are finished with this lab.

To double-check your work, the cell below will rerun all of the autograder tests.

```
In [ ]: grader.check_all()
```

4.13 Submission

Make sure you have run all cells in your notebook in order before running the cell below, so that all images/graphs appear in the output. The cell below will generate a zip file for you to submit. **Please save before exporting!**

```
In [ ]: # Save your notebook first, then run this cell to export your submission.
        grader.export(pdf=False)
```